

GreatSQL TPC-H（Turbo引擎）性能测试报告

GreatSQL TPC-H（Turbo引擎）性能测试报告
(2025年3月5日)

GreatSQL 社区

【文档声明】

GreatSQL 社区提醒您在阅读或使用本文档之前仔细阅读、充分理解本法律声明各条款的内容。如果您阅读或使用本文档，您的阅读或使用行为将被视为对本声明全部内容的认可。您应当通过 GreatSQL 社区网站或 GreatSQL 社区提供的其他授权通道下载、获取本文档，且仅能用于自身的合法合规的业务活动。本文档的内容视为 GreatSQL 社区的保密信息，您应当严格遵守保密义务；未经 GreatSQL 社区事先书面同意，您不得向任何第三方披露本手册内容或提供给任何第三方使用。

未经 GreatSQL 社区事先书面许可，任何单位、公司或个人不得擅自摘抄、翻译、复制本文档内容的部分或全部，不得以任何方式或途径进行替换和宣传。

由于产品版本升级、调整或其他原因，本文档内容有可能变更。GreatSQL 社区保留在没有任何通知或者提示下对本文档的内容进行修改的权利，并在 GreatSQL 社区授权通道中不定期发布更新后的用户文档。您应当实时关注用户文档的版本变更并通过 GreatSQL 社区授权渠道下载、获取最新版的用户文档。

本文档仅作为用户使用 GreatSQL 社区产品及服务的参考性指引。GreatSQL 社区在现有技术的基础上尽最大努力提供相应的介绍及操作指引，但 GreatSQL 社区在此明确声明对本文档内容的准确性、完整性、适用性、可靠性等不作任何明示或暗示的保证。任何单位、公司或个人因为下载、使用或信赖本文档而发生任何差错或经济损失的，GreatSQL 社区不承担任何法律责任。在任何情况下，GreatSQL 社区均不对任何间接性、后果性、惩戒性、偶然性、特殊性或刑罚性的损害，包括用户使用或信赖本文档而遭受的利润损失，承担责任（即使 GreatSQL 社区已被告知该等损失的可能性）。

GreatSQL 社区文档中所有内容，包括但不限于图片、架构设计、页面布局、文字描述，均由 GreatSQL 社区和/或其关联公司依法拥有其知识产权，包括但不限于商标权、专利权、著作权、商业秘密等。非经 GreatSQL 社区和/或其关联公司书面同意，任何人不得擅自使用、修改、复制、公开替换、改变、散布、发行或公开发表 GreatSQL 社区网站、产品程序或内容。此外，未经 GreatSQL 社区事先书面同意，任何人不得为了任何营销、广告、促销或其他目的使用、公布或复制 GreatSQL 社区的名称（包括但不限于单独为或以组合形式包含“GreatSQL 社区”、“GreatSQL”等 GreatSQL 社区和/或其关联公司品牌，上述品牌的附属标志及图案或任何类似公司名称、商号、商标、产品或服务名称、域名、图案标示、标志、标识或通过特定描述使第三方能够识别 GreatSQL 社区和/或其关联公司）。

如若发现本文档存在任何错误，请与 GreatSQL 社区取得直接联系。

GreatSQL社区官网：<https://greatsql.cn>。

1. 概述

本次测试针对GreatSQL数据库Turbo引擎基于标准 TPC-H 场景的测试。

TPC-H（商业智能计算测试）是美国交易处理效能委员会（TPC，TransactionProcessing Performance Council）组织制定的用来模拟决策支持类应用的一个测试集。目前，学术界和工业界普遍采用 TPC-H 来评价决策支持技术方面应用的性能。这种商业测试可以全方位评测系统的整体商业计算综合能力，对厂商的要求更高，同时也具有普遍的商业实用意义，目前在银行信贷分析和信用卡分析、电信运营分析、税收分析、烟草行业决策分析中都有广泛的应用，TPC-H 查询包含八张数据表和 22 条复杂 SQL 查询，大多数查询包含多表联接（JOIN）、子查询和聚合查询等。

GreatSQL数据库是一款**开源免费**数据库，可在普通硬件上满足金融级应用场景，具有**高可用、高性能、高兼容、高安全**等特性，可作为MySQL或Percona Server for MySQL的理想可选替换。

2. 测试环境信息

操作系统	OS: CentOS Linux release 7.9.2009 (Core) 内核: 3.10.0-1160.el7.x86_64
CPU	Intel(R) Xeon(R) Gold 6238 CPU @ 2.10GHz * 4
内存	251G
磁盘	INTEL SSDPE2KE032T8
数据库	GreatSQL 8.0.32-25, Release 25, Revision 79f57097e3f

服务器详细信息

1. 操作系统

```
1  $ cat /etc/os-release
2
3  NAME="CentOS Linux"
4  VERSION="7 (Core)"
5  ID="centos"
6  ID_LIKE="rhel fedora"
7  VERSION_ID="7"
8  PRETTY_NAME="CentOS Linux 7 (Core)"
9  ANSI_COLOR="0;31"
```

```
10 CPE_NAME="cpe:/o:centos:centos:7"
11 HOME_URL="https://www.centos.org/"
12 BUG_REPORT_URL="https://bugs.centos.org/"
13
14 CENTOS_MANTISBT_PROJECT="CentOS-7"
15 CENTOS_MANTISBT_PROJECT_VERSION="7"
16 REDHAT_SUPPORT_PRODUCT="centos"
17 REDHAT_SUPPORT_PRODUCT_VERSION="7"
```

2. CPU

```
1  $ lscpu
2
3  Architecture:          x86_64
4  CPU op-mode(s):        32-bit, 64-bit
5  Byte Order:             Little Endian
6  CPU(s):                 176
7  On-line CPU(s) list:   0-175
8  Thread(s) per core:    2
9  Core(s) per socket:    22
10 Socket(s):              4
11 NUMA node(s):          4
12 Vendor ID:              GenuineIntel
13 CPU family:             6
14 Model:                  85
15 Model name:             Intel(R) Xeon(R) Gold 6238 CPU @ 2.10GHz
16 Stepping:               7
17 CPU MHz:                1000.012
18 CPU max MHz:            3700.0000
19 CPU min MHz:            1000.0000
20 BogomIPS:               4200.00
21 Virtualization:         VT-x
22 L1d cache:              32K
23 L1i cache:              32K
24 L2 cache:               1024K
25 L3 cache:               30976K
26 NUMA node0 CPU(s):
   0,4,8,12,16,20,24,28,32,36,40,44,48,52,56,60,64,68,72,76,80,84,88,92,96,100,10
   4,108,112,116,120,124,128,132,136,140,144,148,152,156,160,164,168,172
27 NUMA node1 CPU(s):
   1,5,9,13,17,21,25,29,33,37,41,45,49,53,57,61,65,69,73,77,81,85,89,93,97,101,10
   5,109,113,117,121,125,129,133,137,141,145,149,153,157,161,165,169,173
28 NUMA node2 CPU(s):
   2,6,10,14,18,22,26,30,34,38,42,46,50,54,58,62,66,70,74,78,82,86,90,94,98,102,1
   06,110,114,118,122,126,130,134,138,142,146,150,154,158,162,166,170,174
```

```
29  NUMA node3 CPU(s):
    3,7,11,15,19,23,27,31,35,39,43,47,51,55,59,63,67,71,75,79,83,87,91,95,99,103,107,111,115,119,123,127,131,135,139,143,147,151,155,159,163,167,171,175
30  Flags:                fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge
    mca cmov pat pse36 clflush dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx
    pdpe1gb rdtscp lm constant_tsc art arch_perfmon pebs bts rep_good nopl
    xtopology nonstop_tsc aperfmperf eagerfpu pni pclmulqdq dtes64 monitor ds_cpl
    vmx smx est tm2 ssse3 sdbg fma cx16 xtpr pdcm pcid dca sse4_1 sse4_2 x2apic
    movbe popcnt tsc_deadline_timer aes xsave avx f16c rdrand lahf_lm abm
    3dnowprefetch epb cat_l3 cdp_l3 invpcid_single intel_ppin intel_pt ssbd mba
    ibrs ibpb stibp ibrs_enhanced tpr_shadow vnmi flexpriority ept vpid fsgsbase
    tsc_adjust bmi1 hle avx2 smep bmi2 erms invpcid rtm cqm mpx rdt_a avx512f
    avx512dq      rdseed adx smap clflushopt clwb avx512cd avx512bw avx512vl
    xsaveopt xsavec xgetbv1 cqm_llc cqm_occup_llc cqm_mbm_total cqm_mbm_local
    dtherm ida arat pln pts pku ospke      avx512_vnni md_clear spec_ctrl
    intel_stibp flush_lld arch_capabilities
```

3. 内存

```
1  $ free -ht
2
3  total      used      free      shared  buff/cache   available
4  Mem:      251G    167G     22G       7.2M        61G       82G
5  Swap:      4.0G     1.1G     2.9G
6  Total:     255G    168G     24G
```

4. 磁盘

磁盘设备型号

```
1  $ nvme list
2  Node              SN                      Model                      Namespace Usage
3  Format              FW  Rev
4  -----
5  /dev/nvme0n1      PHLN018200FD3P2BGN    INTEL SSDPE2KE032T8      1
6  3.20 TB / 3.20 TB  512 B + 0 B    VDV10152
```

磁盘挂载参数、文件系统

```
1  $ df -hT | grep ssd
2  /dev/nvme0n1      xfs          3.0T  1.5T  1.5T  49% /ssd2
```

NVMe SSD设备简单测速

```
1 $ dd oflag=direct if=/dev/zero of=./zero bs=1M count=20480
2
3 20480+0 records in
4 20480+0 records out
5 21474836480 bytes (21 GB) copied, 8.69131 s, 2.5 GB/s
```

 **提示：在下面运行TPC-H测试时，设置了Turbo引擎最大可使用的内存及线程数。**

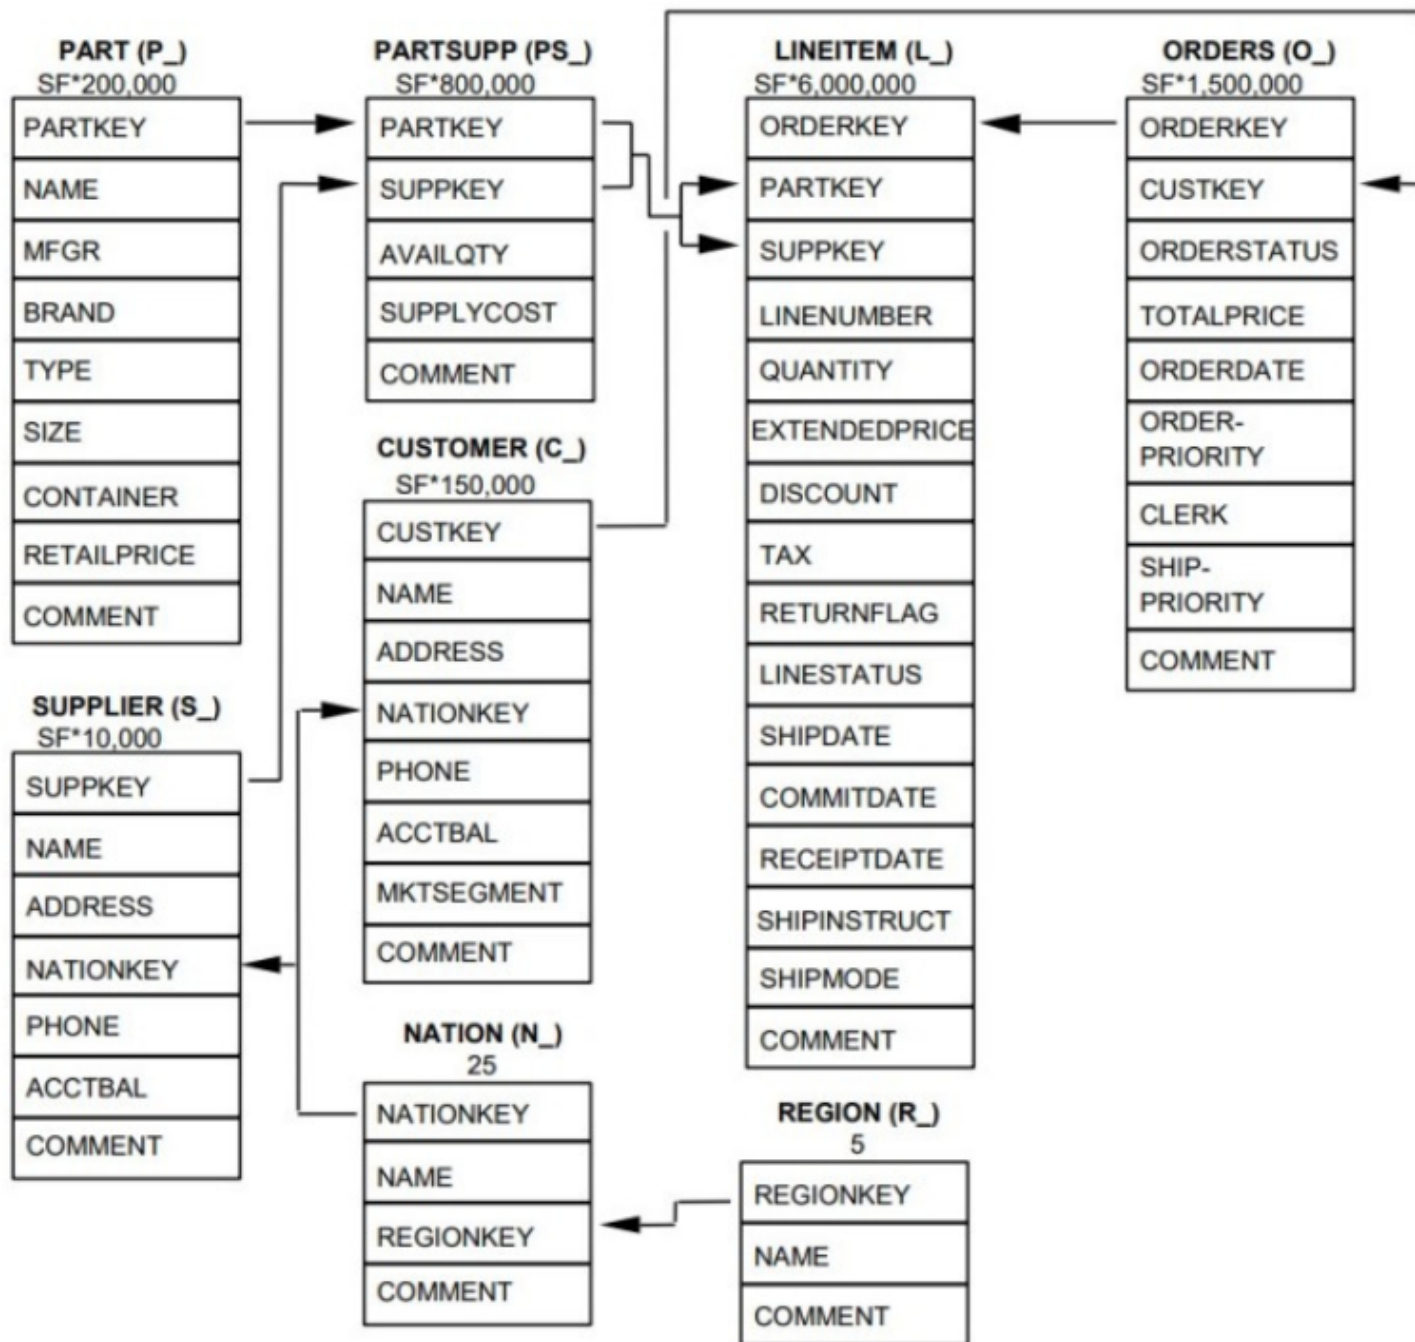
```
1 SET GLOBAL turbo_memory_limit=68719476736;
2 SET GLOBAL turbo_worker_threads=32;
```

3. 测试表结构和数据量

各表数据量对比：

表名	TPC-H SF100数据量	TPC-H SF300数据量	备注
region	5	5	地区信息
nation	25	25	国家表
supplier	1000000	3000000	供应商信息
part	20000000	60000000	零件表
customer	15000000	45000000	消费者表
partsupp	80000000	240000000	配件供应表
orders	150000000	450000000	订单表
lineitem	600037902	1799989091	订单明细表

各表结构关系如下图所示：



4. 测试结果

GreatSQL 8.0.32-27中新增的高性能并行查询引擎Turbo，使得其在实时TPC-H性能测试表现明显优于MySQL社区版、Percona Server MySQL、MariaDB等数据库。


在TPC-H SF100场景下，运行完全部22个TPC-H查询SQL总耗时为 **214.951秒**。在TPC-H SF300场景下，运行完全部22个TPC-H查询SQL总耗时为**2241.448秒**。

每条SQL详细耗时如下：

TPC-H Query (Turbo)	GreatSQL TPC-H SF100 (32C64G) 耗时 (秒)	GreatSQL TPC-H SF300 (32C64G) 耗时 (秒)
Q1	13.839	87.722

Q2	4.684	9.769
Q3	8.956	109.921
Q4	6.618	106.168
Q5	8.673	110.856
Q6	6.723	79.182
Q7	14.645	120.183
Q8	11.669	114.955
Q9	14.845	134.707
Q10	10.359	114.493
Q11	3.022	7.282
Q12	8.092	107.227
Q13	8.988	23.133
Q14	7.878	84.632
Q15	13.051	168.845
Q16	2.934	5.460
Q17	9.779	164.819
Q18	19.197	216.600
Q19	11.355	92.757
Q20	8.114	104.129
Q21	18.073	271.767
Q22	3.457	6.841
总耗时	214.951	2241.448

GreatSQL TPC-H（Turbo）SF100 vs SF300（32C64G）对比示意图如下



5. 测试步骤

5.1 安装 GreatSQL

请参考GreatSQL手册内容：[安装指南](#)，完成GreatSQL安装。

5.2 生成 TPC-H 测试数据

请参考GreatSQL手册内容：[TPC-H性能测试](#)，完成TPC-H工具编译安装。

运行 TPC-H `dbgen` 工具，生成数据文件，一共会生成 8 个表对应的 tbl 数据文件，例如：

```
1  $ ./dbgen -vf -s 100
2  ...
3
4  $ ls -l *tbl
5  -rw-r--r-- 1 root root 2463490271 Sep 26 09:20 customer.tbl
6  -rw-r--r-- 1 root root 79579694556 Sep 26 09:20 lineitem.tbl
7  -rw-r--r-- 1 root root      2224 Sep 26 09:20 nation.tbl
8  -rw-r--r-- 1 root root 17793116301 Sep 26 09:20 orders.tbl
9  -rw-r--r-- 1 root root 12209211160 Sep 26 09:20 partsupp.tbl
10 -rw-r--r-- 1 root root 2453234158 Sep 26 09:20 part.tbl
11 -rw-r--r-- 1 root root      389 Sep 26 09:20 region.tbl
12 -rw-r--r-- 1 root root 142869803 Sep 26 09:20 supplier.tbl
```

也可以参考 [pdbgen.sh](#) 脚本做法，并行生成测试数据。

5.3 创建 TPC-H 测试数据库表并导入数据

参考GreatSQL社区提供的TPC-H数据库表初始化脚本：[tpch-create-table.sql](#)，完成TPC-H测试数据库表创建。

```
1  $ mysql -f < tpch-create-table.sql
2  $ mysqlshow tpch100
3  Database: tpch100
4  +-----+
5  | Tables |
6  +-----+
7  | customer |
8  | lineitem |
9  | nation   |
10 | orders   |
11 | part      |
12 | partsupp |
13 | region   |
14 | revenue0 |
15 | supplier |
16 +-----+
```

利用GreatSQL的 **parallel load data特性** 并行导入TPC-H测试数据。

需要先修改GreatSQL选项 `secure_file_priv` 设置，指向上述 `workdir` 所在目录，重启GreatSQL使之生效。

参考GreatSQL社区提供的并发导入脚本：[load-data-parallel.sh](#)，完成数据导入。

也可以参考 [pload.sh](#) 脚本做法，并行导入数据。

 **提示：运行LOAD DATA导入数据时，可能会在 tmpdir 产生临时文件，因此要保证 tmpdir 有足够的剩余可用磁盘空间。**

5.4 确认Turbo引擎设置

数据导入完成后，在开始运行TPC-H测试前，要先加载Turbo引擎，并修改可用内存及并发线程数等相关设置。

1. 安装Turbo引擎

```
1  INSTALL PLUGIN turbo SONAME 'turbo.so';
```

2. 查看Turbo引擎已安装成功

```

1  greatsql> SELECT * FROM information_schema.PLUGINS WHERE PLUGIN_NAME='turbo'\G
2  ***** 1. row *****
3          PLUGIN_NAME: turbo
4          PLUGIN_VERSION: 0.1
5          PLUGIN_STATUS: ACTIVE
6          PLUGIN_TYPE: QUERY PLAN
7          PLUGIN_TYPE_VERSION: 1.0
8          PLUGIN_LIBRARY: turbo.so
9  PLUGIN_LIBRARY_VERSION: 1.11
10         PLUGIN_AUTHOR: GreatOpenSource
11         PLUGIN_DESCRIPTION: turbo query plan
12         PLUGIN_LICENSE: GPL
13         LOAD_OPTION: ON

```

3. 修改可用内存及并发线程数等相关设置

```

1  greatsql> SET GLOBAL turbo_memory_limit=68719476736;
2  greatsql> SET GLOBAL turbo_worker_threads=32;

```

5.5 执行 TPC-H 测试

执行[GreatSQL社区提供的Turbo引擎专用TPC-H性能测试脚本](#)，完成测试，并记录各个SQL的耗时。

该测试脚本大概工作模式如下：

1. 总共有22个查询SQL，每个查询SQL分别执行。
2. 每个查询SQL先运行2次完成数据预热。
3. 每个SQL再执行3次，每次执行SQL都会记录其起止时间，及其耗时，如下面例所示：

```

1  [2023-09-27 01:38:45] BEGIN RUN TPC-H Q1 1 times
2  [2023-09-27 01:38:46] TPC-H Q1 END, COST: 1.301s
3
4
5  [2023-09-27 01:38:46] BEGIN RUN TPC-H Q1 2 times
6  [2023-09-27 01:38:47] TPC-H Q1 END, COST: 0.787s

```

上述结果中的 COST: 1.301s ，即为本SQL的运行耗时：1.301秒。

4. 继续执行下一个查询SQL，直至22个查询SQL全部执行完毕。

可以参考自动化执行脚本 [run-tpch.sh](#) 的做法，修改几个参数后即可自动执行。

 提示：在运行 `tpch_queries_11.sql` 这个SQL脚本时，需要根据数据量大小调整第17-20行相关的参数。例如当测试数据量是SF100时，调整成如下

```
1  SELECT /* SET_VAR(turbo_enable=ON) SET_VAR(turbo_cost_threshold=0) */ /*+  
   Q11 */  
2      ps_partkey,  
3      sum(ps_supplycost * ps_availqty) AS value  
4  FROM  
5      partsupp,  
6      supplier,  
7      nation  
8  WHERE  
9      ps_suppkey = s_suppkey  
10     AND s_nationkey = n_nationkey  
11     AND n_name = 'GERMANY'  
12 GROUP BY  
13     ps_partkey  
14 HAVING  
15     sum(ps_supplycost * ps_availqty) > (  
16         SELECT  
17             /* sum(ps_supplycost * ps_availqty) * 0.0001000000 /* SF1 */  
18             /* sum(ps_supplycost * ps_availqty) * 0.0000100000 /* SF10 */  
19             sum(ps_supplycost * ps_availqty) * 0.0000010000 /* SF100 */  
20             /* sum(ps_supplycost * ps_availqty) * 0.0000001000 /* SF1000 */  
21         FROM  
22             partsupp,  
23             supplier,  
24             nation  
25         WHERE  
26             ps_suppkey = s_suppkey  
27             AND s_nationkey = n_nationkey  
28             AND n_name = 'GERMANY')  
29 ORDER BY  
30     value DESC;
```

6. 附录

6.1 创建测试表DDL

```
1  -- DROP DATABASE IF EXISTS tpch;  
2  -- CREATE DATABASE IF NOT EXISTS tpch DEFAULT CHARACTER SET latin1;
```



```

46         ps_supplycost decimal(15,2) not null,
47         ps_comment      varchar(199) not null,
48         primary key(ps_partkey,ps_suppkey),
49         key partsupp_fk1 (ps_suppkey),
50         key partsupp_fk2 (ps_partkey) )
secondary_engine = rapid;
51
52
53 drop table if exists customer;
54 create table customer ( c_custkey      integer not null,
55                        c_name          varchar(25) not null,
56                        c_address       varchar(40) not null,
57                        c_nationkey     integer not null,
58                        c_phone         char(15) not null,
59                        c_acctbal       decimal(15,2) not null,
60                        c_mktsegment   char(10) not null,
61                        c_comment      varchar(117) not null,
62                        primary key(c_custkey),
63                        key customer_fk1 (c_nationkey) )
secondary_engine = rapid;
64
65 drop table if exists orders;
66 create table orders ( o_orderkey      integer not null,
67                     o_custkey        integer not null,
68                     o_orderstatus    char(1) not null,
69                     o_totalprice     decimal(15,2) not null,
70                     o_orderdate      date not null,
71                     o_orderpriority char(15) not null,
72                     o_clerk          char(15) not null,
73                     o_shippriority  integer not null,
74                     o_comment       varchar(79) not null,
75                     primary key(o_orderkey),
76                     key orders_fk1 (o_custkey) ) secondary_engine
= rapid;
77
78 drop table if exists lineitem;
79 create table lineitem ( l_orderkey    integer not null,
80                      l_partkey      integer not null,
81                      l_suppkey      integer not null,
82                      l_linenum      integer not null,
83                      l_quantity     decimal(15,2) not null,
84                      l_extendedprice decimal(15,2) not null,
85                      l_discount     decimal(15,2) not null,
86                      l_tax          decimal(15,2) not null,
87                      l_returnflag   char(1) not null,
88                      l_linestatus   char(1) not null,
89                      l_shipdate     date not null,

```

```

90         l_commitdate date not null,
91         l_receiptdate date not null,
92         l_shipinstruct char(25) not null,
93         l_shipmode      char(10) not null,
94         l_comment        varchar(44) not null,
95         primary key(l_orderkey,l_linenum),
96         key lineitem_fk1 (l_orderkey) ,
97         key lineitem_fk2 (l_partkey,l_suppkey) )
secondary_engine = rapid;

```

6.2 22条TPC-H测试SQL

```

1  -- tpch_queries_1.sql
2  SELECT /*+ SET_VAR(use_secondary_engine=1)
   SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q1 */
3      l_returnflag,
4      l_linestatus,
5      sum(l_quantity) AS sum_qty,
6      sum(l_extendedprice) AS sum_base_price,
7      sum(l_extendedprice * (1 - l_discount)) AS sum_disc_price,
8      sum(l_extendedprice * (1 - l_discount) * (1 + l_tax)) AS sum_charge,
9      avg(l_quantity) AS avg_qty,
10     avg(l_extendedprice) AS avg_price,
11     avg(l_discount) AS avg_disc,
12     count(*) AS count_order
13 FROM
14     lineitem
15 WHERE
16     l_shipdate <= CAST('1998-09-02' AS date)
17 GROUP BY
18     l_returnflag,
19     l_linestatus
20 ORDER BY
21     l_returnflag,
22     l_linestatus;
23
24
25
26
27 -- tpch_queries_2.sql
28 SELECT /*+ SET_VAR(use_secondary_engine=1)
   SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q2 */
29     s_acctbal,
30     s_name,
31     n_name,

```

```

32     p_partkey,
33     p_mfgr,
34     s_address,
35     s_phone,
36     s_comment
37 FROM
38     part,
39     supplier,
40     partsupp,
41     nation,
42     region
43 WHERE
44     p_partkey = ps_partkey
45     AND s_suppkey = ps_suppkey
46     AND p_size = 15
47     AND p_type LIKE '%BRASS'
48     AND s_nationkey = n_nationkey
49     AND n_regionkey = r_regionkey
50     AND r_name = 'EUROPE'
51     AND ps_supplycost = (
52         SELECT
53             min(ps_supplycost)
54         FROM
55             partsupp,
56             supplier,
57             nation,
58             region
59         WHERE
60             p_partkey = ps_partkey
61             AND s_suppkey = ps_suppkey
62             AND s_nationkey = n_nationkey
63             AND n_regionkey = r_regionkey
64             AND r_name = 'EUROPE')
65 ORDER BY
66     s_acctbal DESC,
67     n_name,
68     s_name,
69     p_partkey
70 LIMIT 100;
71
72
73
74
75 -- tpch_queries_3.sql
76 SELECT /*+ SET_VAR(use_secondary_engine=1)
77         SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q3 */
78     l_orderkey,

```

```

78         sum(l_extendedprice * (1 - l_discount)) AS revenue,
79         o_orderdate,
80         o_shippriority
81 FROM
82     customer,
83     orders,
84     lineitem
85 WHERE
86     c_mktsegment = 'BUILDING'
87     AND c_custkey = o_custkey
88     AND l_orderkey = o_orderkey
89     AND o_orderdate < CAST('1995-03-15' AS date)
90     AND l_shipdate > CAST('1995-03-15' AS date)
91 GROUP BY
92     l_orderkey,
93     o_orderdate,
94     o_shippriority
95 ORDER BY
96     revenue DESC,
97     o_orderdate
98 LIMIT 10;
99
100
101
102
103 -- tpch_queries_4.sql
104 SELECT /*+ SET_VAR(use_secondary_engine=1)
105         SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q4 */
106         o_orderpriority,
107         count(*) AS order_count
108 FROM
109     orders
110 WHERE
111     o_orderdate >= CAST('1993-07-01' AS date)
112     AND o_orderdate < CAST('1993-10-01' AS date)
113     AND EXISTS (
114         SELECT
115             *
116         FROM
117             lineitem
118         WHERE
119             l_orderkey = o_orderkey
120             AND l_commitdate < l_receiptdate)
121 GROUP BY
122     o_orderpriority
123 ORDER BY
124     o_orderpriority;

```



```
124
125
126
127
128 -- tpch_queries_5.sql
129 SELECT /*+ SET_VAR(use_secondary_engine=1)
      SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q5 */
130     n_name,
131     sum(l_extendedprice * (1 - l_discount)) AS revenue
132 FROM
133     customer,
134     orders,
135     lineitem,
136     supplier,
137     nation,
138     region
139 WHERE
140     c_custkey = o_custkey
141     AND l_orderkey = o_orderkey
142     AND l_suppkey = s_suppkey
143     AND c_nationkey = s_nationkey
144     AND s_nationkey = n_nationkey
145     AND n_regionkey = r_regionkey
146     AND r_name = 'ASIA'
147     AND o_orderdate >= CAST('1994-01-01' AS date)
148     AND o_orderdate < CAST('1995-01-01' AS date)
149 GROUP BY
150     n_name
151 ORDER BY
152     revenue DESC;
153
154
155
156
157 -- tpch_queries_6.sql
158 SELECT /*+ SET_VAR(use_secondary_engine=1)
      SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q6 */
159     sum(l_extendedprice * l_discount) AS revenue
160 FROM
161     lineitem
162 WHERE
163     l_shipdate >= CAST('1994-01-01' AS date)
164     AND l_shipdate < CAST('1995-01-01' AS date)
165     AND l_discount BETWEEN 0.05
166     AND 0.07
167     AND l_quantity < 24;
168
```

```

169
170
171
172 -- tpch_queries_7.sql
173 SELECT /*+ SET_VAR(use_secondary_engine=1)
      SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q7 */
174     supp_nation,
175     cust_nation,
176     l_year,
177     sum(volume) AS revenue
178 FROM (
179     SELECT
180         n1.n_name AS supp_nation,
181         n2.n_name AS cust_nation,
182         extract(year FROM l_shipdate) AS l_year,
183         l_extendedprice * (1 - l_discount) AS volume
184     FROM
185         supplier,
186         lineitem,
187         orders,
188         customer,
189         nation n1,
190         nation n2
191     WHERE
192         s_suppkey = l_suppkey
193         AND o_orderkey = l_orderkey
194         AND c_custkey = o_custkey
195         AND s_nationkey = n1.n_nationkey
196         AND c_nationkey = n2.n_nationkey
197         AND ((n1.n_name = 'FRANCE'
198             AND n2.n_name = 'GERMANY')
199             OR (n1.n_name = 'GERMANY'
200             AND n2.n_name = 'FRANCE'))
201         AND l_shipdate BETWEEN CAST('1995-01-01' AS date)
202             AND CAST('1996-12-31' AS date)) AS shipping
203 GROUP BY
204     supp_nation,
205     cust_nation,
206     l_year
207 ORDER BY
208     supp_nation,
209     cust_nation,
210     l_year;
211
212
213
214

```

```

215  -- tpch_queries_8.sql
216  SELECT /*+ SET_VAR(use_secondary_engine=1)
      SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q8 */
217      o_year,
218      sum(
219          CASE WHEN nation = 'BRAZIL' THEN
220              volume
221          ELSE
222              0
223          END) / sum(volume) AS mkt_share
224  FROM (
225      SELECT
226          extract(year FROM o_orderdate) AS o_year,
227          l_extendedprice * (1 - l_discount) AS volume,
228          n2.n_name AS nation
229      FROM
230          part,
231          supplier,
232          lineitem,
233          orders,
234          customer,
235          nation n1,
236          nation n2,
237          region
238      WHERE
239          p_partkey = l_partkey
240          AND s_suppkey = l_suppkey
241          AND l_orderkey = o_orderkey
242          AND o_custkey = c_custkey
243          AND c_nationkey = n1.n_nationkey
244          AND n1.n_regionkey = r_regionkey
245          AND r_name = 'AMERICA'
246          AND s_nationkey = n2.n_nationkey
247          AND o_orderdate BETWEEN CAST('1995-01-01' AS date)
248              AND CAST('1996-12-31' AS date)
249          AND p_type = 'ECONOMY ANODIZED STEEL') AS all_nations
250  GROUP BY
251      o_year
252  ORDER BY
253      o_year;
254
255
256
257
258  -- tpch_queries_9.sql
259  SELECT /*+ SET_VAR(use_secondary_engine=1)
      SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q9 */

```

```

260     nation,
261     o_year,
262     sum(amount) AS sum_profit
263 FROM (
264     SELECT
265         n_name AS nation,
266         extract(year FROM o_orderdate) AS o_year,
267         l_extendedprice * (1 - l_discount) - ps_supplycost * l_quantity AS
amount
268     FROM
269         part,
270         supplier,
271         lineitem,
272         partsupp,
273         orders,
274         nation
275     WHERE
276         s_suppkey = l_suppkey
277         AND ps_suppkey = l_suppkey
278         AND ps_partkey = l_partkey
279         AND p_partkey = l_partkey
280         AND o_orderkey = l_orderkey
281         AND s_nationkey = n_nationkey
282         AND p_name LIKE '%green%') AS profit
283 GROUP BY
284     nation,
285     o_year
286 ORDER BY
287     nation,
288     o_year DESC;
289
290
291
292
293 -- tpch_queries_10.sql
294 SELECT /*+ SET_VAR(use_secondary_engine=1)
SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q10 */
295     c_custkey,
296     c_name,
297     sum(l_extendedprice * (1 - l_discount)) AS revenue,
298     c_acctbal,
299     n_name,
300     c_address,
301     c_phone,
302     c_comment
303 FROM
304     customer,

```

```

305     orders,
306     lineitem,
307     nation
308 WHERE
309     c_custkey = o_custkey
310     AND l_orderkey = o_orderkey
311     AND o_orderdate >= CAST('1993-10-01' AS date)
312     AND o_orderdate < CAST('1994-01-01' AS date)
313     AND l_returnflag = 'R'
314     AND c_nationkey = n_nationkey
315 GROUP BY
316     c_custkey,
317     c_name,
318     c_acctbal,
319     c_phone,
320     n_name,
321     c_address,
322     c_comment
323 ORDER BY
324     revenue DESC
325 LIMIT 20;
326
327
328
329
330 -- tpch_queries_11.sql
331 SELECT /*+ SET_VAR(use_secondary_engine=1)
332        SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q11 */
333     ps_partkey,
334     sum(ps_supplycost * ps_availqty) AS value
335 FROM
336     partsupp,
337     supplier,
338     nation
339 WHERE
340     ps_suppkey = s_suppkey
341     AND s_nationkey = n_nationkey
342     AND n_name = 'GERMANY'
343 GROUP BY
344     ps_partkey
345 HAVING
346     sum(ps_supplycost * ps_availqty) > (
347         SELECT
348             sum(ps_supplycost * ps_availqty) * 0.0001000000 /* SF1 */
349             /* sum(ps_supplycost * ps_availqty) * 0.0000100000 /* SF10 */
350             /* sum(ps_supplycost * ps_availqty) * 0.0000010000 /* SF100 */

```

```

350          /* sum(ps_supplycost * ps_availqty) * 0.0000001000 /* SF1000 */

351      FROM
352          partsupp,
353          supplier,
354          nation
355      WHERE
356          ps_suppkey = s_suppkey
357          AND s_nationkey = n_nationkey
358          AND n_name = 'GERMANY')
359  ORDER BY
360      value DESC;
361
362
363
364
365  -- tpch_queries_12.sql
366  SELECT /*+ SET_VAR(use_secondary_engine=1)
367  SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q12 */
368      l_shipmode,
369      sum(
370          CASE WHEN o_orderpriority = '1-URGENT'
371              OR o_orderpriority = '2-HIGH' THEN
372              1
373          ELSE
374              0
375          END) AS high_line_count,
376      sum(
377          CASE WHEN o_orderpriority <> '1-URGENT'
378              AND o_orderpriority <> '2-HIGH' THEN
379              1
380          ELSE
381              0
382          END) AS low_line_count
383  FROM
384      orders,
385      lineitem
386  WHERE
387      o_orderkey = l_orderkey
388      AND l_shipmode IN ('MAIL', 'SHIP')
389      AND l_commitdate < l_receiptdate
390      AND l_shipdate < l_commitdate
391      AND l_receiptdate >= CAST('1994-01-01' AS date)
392      AND l_receiptdate < CAST('1995-01-01' AS date)
393  GROUP BY
394      l_shipmode
395  ORDER BY

```

```

395         l_shipmode;
396
397
398
399
400 -- tpch_queries_13.sql
401 SELECT /*+ SET_VAR(use_secondary_engine=1)
         SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q13 */
402         c_count,
403         count(*) AS custdist
404 FROM (
405     SELECT
406         c_custkey,
407         count(o_orderkey)
408     FROM
409         customer
410     LEFT OUTER JOIN orders ON c_custkey = o_custkey
411     AND o_comment NOT LIKE '%special%requests%'
412 GROUP BY
413     c_custkey) AS c_orders (c_custkey,
414     c_count)
415 GROUP BY
416     c_count
417 ORDER BY
418     custdist DESC,
419     c_count DESC;
420
421
422
423
424 -- tpch_queries_14.sql
425 SELECT /*+ SET_VAR(use_secondary_engine=1)
         SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q14 */
426     100.00 * sum(
427         CASE WHEN p_type LIKE 'PROMO%' THEN
428             l_extendedprice * (1 - l_discount)
429         ELSE
430             0
431         END) / sum(l_extendedprice * (1 - l_discount)) AS promo_revenue
432 FROM
433     lineitem,
434     part
435 WHERE
436     l_partkey = p_partkey
437     AND l_shipdate >= date '1995-09-01'
438     AND l_shipdate < CAST('1995-10-01' AS date);
439

```

```

440
441
442
443 -- tpch_queries_15.sql
444 SELECT /*+ SET_VAR(use_secondary_engine=1)
      SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q15 */
445     s_suppkey,
446     s_name,
447     s_address,
448     s_phone,
449     total_revenue
450 FROM
451     supplier,
452     (
453         SELECT
454             l_suppkey AS supplier_no,
455             sum(l_extendedprice * (1 - l_discount)) AS total_revenue
456         FROM
457             lineitem
458         WHERE
459             l_shipdate >= CAST('1996-01-01' AS date)
460             AND l_shipdate < CAST('1996-04-01' AS date)
461         GROUP BY
462             supplier_no) revenue0
463 WHERE
464     s_suppkey = supplier_no
465     AND total_revenue = (
466         SELECT
467             max(total_revenue)
468         FROM (
469             SELECT
470                 l_suppkey AS supplier_no,
471                 sum(l_extendedprice * (1 - l_discount)) AS total_revenue
472             FROM
473                 lineitem
474             WHERE
475                 l_shipdate >= CAST('1996-01-01' AS date)
476                 AND l_shipdate < CAST('1996-04-01' AS date)
477             GROUP BY
478                 supplier_no) revenue1)
479 ORDER BY
480     s_suppkey;
481
482
483
484
485 -- tpch_queries_16.sql

```



```

486 SELECT /*+ SET_VAR(use_secondary_engine=1)
      SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q16 */
487     p_brand,
488     p_type,
489     p_size,
490     count(DISTINCT ps_suppkey) AS supplier_cnt
491 FROM
492     partsupp,
493     part
494 WHERE
495     p_partkey = ps_partkey
496     AND p_brand <> 'Brand#45'
497     AND p_type NOT LIKE 'MEDIUM POLISHED%'
498     AND p_size IN (49, 14, 23, 45, 19, 3, 36, 9)
499     AND ps_suppkey NOT IN (
500         SELECT
501             s_suppkey
502         FROM
503             supplier
504         WHERE
505             s_comment LIKE '%Customer%Complaints%')
506 GROUP BY
507     p_brand,
508     p_type,
509     p_size
510 ORDER BY
511     supplier_cnt DESC,
512     p_brand,
513     p_type,
514     p_size;
515
516
517
518
519 -- tpch_queries_17.sql
520 SELECT /*+ SET_VAR(use_secondary_engine=1)
      SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q17 */
521     sum(l_extendedprice) / 7.0 AS avg_yearly
522 FROM
523     lineitem,
524     part
525 WHERE
526     p_partkey = l_partkey
527     AND p_brand = 'Brand#23'
528     AND p_container = 'MED BOX'
529     AND l_quantity < (
530         SELECT

```

```

531         0.2 * avg(l_quantity)
532     FROM
533         lineitem
534     WHERE
535         l_partkey = p_partkey);
536
537
538
539
540 -- tpch_queries_18.sql
541 SELECT /*+ SET_VAR(use_secondary_engine=1)
542 SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q18 */
543     c_name,
544     c_custkey,
545     o_orderkey,
546     o_orderdate,
547     o_totalprice,
548     sum(l_quantity)
549 FROM
550     customer,
551     orders,
552     lineitem
553 WHERE
554     o_orderkey IN (
555         SELECT
556             l_orderkey
557         FROM
558             lineitem
559         GROUP BY
560             l_orderkey
561         HAVING
562             sum(l_quantity) > 300)
563     AND c_custkey = o_custkey
564     AND o_orderkey = l_orderkey
565 GROUP BY
566     c_name,
567     c_custkey,
568     o_orderkey,
569     o_orderdate,
570     o_totalprice
571 ORDER BY
572     o_totalprice DESC,
573     o_orderdate
574 LIMIT 100;
575
576

```

```

577
578 -- tpch_queries_19.sql
579 SELECT /*+ SET_VAR(use_secondary_engine=1)
      SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q19 */
580     sum(l_extendedprice * (1 - l_discount)) AS revenue
581 FROM
582     lineitem,
583     part
584 WHERE (p_partkey = l_partkey
585        AND p_brand = 'Brand#12'
586        AND p_container IN ('SM CASE', 'SM BOX', 'SM PACK', 'SM PKG')
587        AND l_quantity >= 1
588        AND l_quantity <= 1 + 10
589        AND p_size BETWEEN 1 AND 5
590        AND l_shipmode IN ('AIR', 'AIR REG')
591        AND l_shipinstruct = 'DELIVER IN PERSON')
592 OR (p_partkey = l_partkey
593     AND p_brand = 'Brand#23'
594     AND p_container IN ('MED BAG', 'MED BOX', 'MED PKG', 'MED PACK')
595     AND l_quantity >= 10
596     AND l_quantity <= 10 + 10
597     AND p_size BETWEEN 1 AND 10
598     AND l_shipmode IN ('AIR', 'AIR REG')
599     AND l_shipinstruct = 'DELIVER IN PERSON')
600 OR (p_partkey = l_partkey
601     AND p_brand = 'Brand#34'
602     AND p_container IN ('LG CASE', 'LG BOX', 'LG PACK', 'LG PKG')
603     AND l_quantity >= 20
604     AND l_quantity <= 20 + 10
605     AND p_size BETWEEN 1 AND 15
606     AND l_shipmode IN ('AIR', 'AIR REG')
607     AND l_shipinstruct = 'DELIVER IN PERSON');
608
609
610
611
612 -- tpch_queries_20.sql
613 SELECT /*+ SET_VAR(use_secondary_engine=1)
      SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q20 */
614     s_name,
615     s_address
616 FROM
617     supplier,
618     nation
619 WHERE
620     s_suppkey IN (
621         SELECT

```

```

622         ps_suppkey
623     FROM
624         partsupp
625     WHERE
626         ps_partkey IN (
627             SELECT
628                 p_partkey
629             FROM
630                 part
631             WHERE
632                 p_name LIKE 'forest%')
633         AND ps_availqty > (
634             SELECT
635                 0.5 * sum(l_quantity)
636             FROM
637                 lineitem
638             WHERE
639                 l_partkey = ps_partkey
640                 AND l_suppkey = ps_suppkey
641                 AND l_shipdate >= CAST('1994-01-01' AS date)
642                 AND l_shipdate < CAST('1995-01-01' AS date)))
643         AND s_nationkey = n_nationkey
644         AND n_name = 'CANADA'
645     ORDER BY
646         s_name;
647
648
649
650
651 -- tpch_queries_21.sql
652 SELECT /*+ SET_VAR(use_secondary_engine=1)
653        SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q21 */
654     s_name,
655     count(*) AS numwait
656 FROM
657     supplier,
658     lineitem l1,
659     orders,
660     nation
661 WHERE
662     s_suppkey = l1.l_suppkey
663     AND o_orderkey = l1.l_orderkey
664     AND o_orderstatus = 'F'
665     AND l1.l_receiptdate > l1.l_commitdate
666     AND EXISTS (
667         SELECT
668             *

```

```

668         FROM
669             lineitem l2
670         WHERE
671             l2.l_orderkey = l1.l_orderkey
672             AND l2.l_suppkey <> l1.l_suppkey)
673     AND NOT EXISTS (
674         SELECT
675             *
676         FROM
677             lineitem l3
678         WHERE
679             l3.l_orderkey = l1.l_orderkey
680             AND l3.l_suppkey <> l1.l_suppkey
681             AND l3.l_receiptdate > l3.l_commitdate)
682     AND s_nationkey = n_nationkey
683     AND n_name = 'SAUDI ARABIA'
684 GROUP BY
685     s_name
686 ORDER BY
687     numwait DESC,
688     s_name
689 LIMIT 100;
690
691
692
693
694 -- tpch_queries_22.sql
695 SELECT /*+ SET_VAR(use_secondary_engine=1)
696         SET_VAR(secondary_engine_cost_threshold=0) */ /*+ Q22 */
697     cntrycode,
698     count(*) AS numcust,
699     sum(c_acctbal) AS totacctbal
700 FROM (
701     SELECT
702         substring(c_phone FROM 1 FOR 2) AS cntrycode,
703         c_acctbal
704     FROM
705         customer
706     WHERE
707         substring(c_phone FROM 1 FOR 2) IN ('13', '31', '23', '29', '30',
708         '18', '17')
709         AND c_acctbal > (
710             SELECT
711                 avg(c_acctbal)
712             FROM
713                 customer
714             WHERE

```

```
713             c_acctbal > 0.00
714             AND substring(c_phone FROM 1 FOR 2) IN ('13', '31', '23',
              '29', '30', '18', '17'))
715             AND NOT EXISTS (
716                 SELECT
717                     *
718                 FROM
719                     orders
720                 WHERE
721                     o_custkey = c_custkey)) AS custsale
722 GROUP BY
723     centrycode
724 ORDER BY
725     centrycode;
```

6.3 参考资料

- TPC-H官网 <http://www.tpc.org/tpch>
- GreatSQL安装指南: <https://greatsql.cn/docs/4-install-guide/0-install-guide.html>
- TPC-H性能测试指南: <https://greatsql.cn/docs/10-optimize/3-2-benchmark-tpch.html>
- TPC-H性能测试工具包: <https://gitee.com/GreatSQL/tpch>